

Reflections About Persistency



A Very Technical Journey

What's here to learn?

- ❖ Criteria for choosing a persistence technology
- ❖ Reasons pro and contra technology abstraction
- ❖ Good judgement for using transaction semantics
- ❖ Good judgement for using caching

What's here to see?

- ❖ Various APIs
 - ❖ JPA, Criteria API,
 - ❖ JOOQ, QueryDSL,
 - ❖ Spring Data
- ❖ We'll discuss a case study
- ❖ and that's where the fun starts...

Our Technology Stack

- ❖ Java SDK 8
- ❖ Spring 4.0 with Spring-Data and Hibernate
JPA 2.1
- ❖ MySQL
- ❖ MongoDB
- ❖ EhCache

If you'd have to program this...

- ❖ Swiss Trusted Poker Site
- ❖ Semi-private internet poker for closed groups
- ❖ Very high stakes, still thousands of players worldwide
- ❖ Extremely reliable
AND blazing fast
- ❖ All bets must be perfectly
consistent (it's loads of money)



Reflecting on the problem

- ❖ What would you choose?
- ❖ What technology?
 - ❖ SQL or NoSQL?
- ❖ What deployment model?
 - ❖ Master-Slave, other HA options
- ❖ What transaction semantics?
 - ❖ CAP or ACID?

Varying Requirements

- ❖ Different Use Cases may require different runtime characteristics
- ❖ Not always, one size fits it all
- ❖ Modern architectures employ different technologies
- ❖ Keep your architectural mind open
- ❖ Mostly ok to consider JPA, the standard, first

JPA is the Standard

- ❖ Java Persistence API 2.1 now the JEE Standard
- ❖ inspired by Toplink and Hibernate
- ❖ Replaced the EJB 2.1 Entity Beans long ago.
- ❖ Now, what is JPA - in a nutshell?

JPA in a Nutshell

- ❖ ORM: Object-relational mapping
- ❖ Map Classes to Tables
- ❖ Map primitive fields to columns
- ❖ Map references to foreign keys
- ❖ Map inheritance in three different ways
- ❖ Map JPQL to SQL

API: the EntityManager

```
@Component
@Transactional
public class ProductRepository {

    @PersistenceContext
    private EntityManager entityManager;

    public void save ( Product product ) { entityManager.persist( product ); }

    public Product findById(String id) {
        TypedQuery<Product> query = entityManager.createQuery(
            "from Product where id = :id", Product.class);

        query.setParameter("id", id );

        return query.getSingleResult();
    }
}
```


Mapping fields

```
▶ @Entity(name = "Contract")  
  @Table(name = "CONTRACTS")  
  public class Contract extends SecureResource {  
  
      private String contractId;  
  
      @Enumerated(EnumType.STRING)  
      private ContractType contractType;  
  
      private Contract() {  
      } // for JPA
```



Mapping references

```
@Entity
@Inheritance(strategy = InheritanceType.JOINED)
@Table(name = "SECURE_RESOURCES")
public abstract class SecureResource extends BaseEntity {

    @ManyToOne
    private User owner;

    @ManyToOne
    private Tenant tenant;
```


The big picture



```
public class Customer {  
  
    @Id @GeneratedValue Long id;  
  
    @OneToOne  
    @JoinColumn (name="address")  
    Address getAddress() {...}  
    ...  
};
```

Customer.java

```
public class Address {  
  
    @Id @GeneratedValue Long id;  
  
    String getStreet () {...}  
    String getCountry () {...}  
    ...  
};
```

Address.java

Lazy vs Eager

- ❖ Not the entire graph will be read from DB
- ❖ References may be lazily loaded.
- ❖ So may be BLOBs
- ❖ Only accessing lazy fields will retrieve referenced entities
- ❖ LazyInitializationException if outside of transactional context.


```
@Test(expected = LazyInitializationException.class)
public void lazyInitializationDemo () {

    List<Book> books = bookService.findByTitle("Myth Busting in SWE");

    String firstName = books.get(0).getAuthor().getFirstName();

    Assert.fail("Shouldn't get here to assert on " + firstName);
}
}
```

Gavin King's nightmare

Lazy-loaded references can't be accessed outside of the transaction

ACID Transactions

- ❖ Atomic: all or nothing to succeed
- ❖ Consistent: obey constraints
- ❖ Isolated: from each other
- ❖ Durable: on the disk



Transactions with Spring

- ❖ JPA can delegate Tx-Handling to the container
- ❖ Transactions are handled at the method entry and exit through AOP
- ❖ Exceptions may cause roll-back
- ❖ @Transactional Annotation defines attributes

Rollback Strategies

- ❖ Runtime Exceptions cause a roll-back
- ❖ Checked Exceptions don't
- ❖ Exceptions to the rule by
 - ❖ `@Transactional (rollbackOn=Class[])`
 - ❖ `@Transactional (dontRollbackOn=Class[])`

Propagation

- ❖ With Spring very similar to JEE
- ❖ REQUIRED: create tx if none available yet
- ❖ REQUIRES_NEW: create tx, suspend existing
- ❖ MANDATORY: fail if none available
- ❖ NOT_SUPPORTED: suspend, if existing
- ❖ SUPPORTS: execute in whatever is available
- ❖ NEVER: fail if tx available

How does a method know,
whether there's a transaction?

Saving Audit Logs

- ❖ Outer Method: REQUIRED
- ❖ Inner Method: REQUIRES_NEW
- ❖ Outer Method:
 - ❖ When exception occurs:
 - ❖ outer context rolls back
 - ❖ inner context still committed.

Caching

- ❖ Client or server-side caching?
- ❖ Single point of access? Otherwise outdated?
- ❖ Evict entry when writing
- ❖ Evict to avoid OutOfMemoryException
- ❖ LRU (last recently used)
- ❖ Life time expiration
- ❖ May be difficult to test

The Caching Aspect

```
@CacheEvict(value = "projects", key = "#project.name")  
public void removeProject ( Project project ) {  
    repo.delete( project );  
}
```

```
@Transactional  
@CacheEvict(value = "projects", key = "#project.name")  
public void createProject ( Project project ) { repo.save ( project ); }
```

```
@Cacheable( value = "projects", key = "#name")  
@Transactional  
public Project findProject ( String name ) {  
  
    LOGGER.info("accessing DB...");  
}
```

❖ run CacheTests.java and see the result

To cache or not to cache?

- ❖ Keep the results of expensive operations in memory, but where exactly?
- ❖ Consider eviction, single point of access
- ❖ Consider security
- ❖ Consider testability
- ❖ JCache in Java 9 (hopefully) and Spring 4.1

RDBMS: different APIs

- ❖ JPA with TypedQuery from Strings
- ❖ JPA with type-safe Criteria API
- ❖ QueryDSL with generated code
- ❖ JOOQ with generated Code
- ❖ See demo test classes in
`org.smurve.hsr2014.apis`


```

/**
 * Select all book titles starting with "My" or "Your"
 * @param ctx the DSL Context to use
 * @return the resulting list of records
 */
private Result<Record> selectSomeRecords(DSLContext ctx) {
    return ctx.select().from(BOOK).join(AUTHOR).on(BOOK.AUTHOR_ID.equal(AUTHOR.ID))
        .where(
            BOOK.AUTHOR_ID.equal(10L))
        .and(
            BOOK.TITLE.startsWith("My")
                .or(BOOK.TITLE.startsWith("Your")))
        .and(AUTHOR.LASTNAME.contains("irsch"))
        .fetch();
}

```

JOOQ is JOOL

Use for complex reporting queries

Low Level APIs

- ❖ Low-level APIs typically need be integrated to have transactions managed by a larger framework
- ❖ In the exercises: JooqSpringDemo


```
// custQuery is an IEnumerable<IGrouping<string, Customer>>
var custQuery =
    from cust in customers
    group cust by cust.City into custGroup
    where custGroup.Count() > 2
    orderby custGroup.Key
    select custGroup;
```

OK, ok, LINQ is cooler



Choosing Persistence

...is not so easy!

RDBMS/ORM/JPA

- ❖ Very well established technology
- ❖ Transactions: ACID
- ❖ JPA: Standardised abstraction from underlying product
- ❖ A relational System is good for things that are significantly related and form small graphs.

Non-Trivial Tx

- ❖ Exceptions and Rollbacks
- ❖ Propagation Semantics for Audit Requirements
- ❖ One example in the exercises (please find out how it works):

```
MultiTenantSecurityTest.  
    test_audit_records_with_Exceptions()
```


ORM: The Abstraction Promise

- ❖ Hide the details from the developer
- ❖ Everything is OO
- ❖ loose coupling
- ❖ make DB replaceable
- ❖ and what's the result?

ORM Gone Bad: The Abstraction Problem

- ❖ Start easily without thinking, think later and thus a lot harder
- ❖ Many non-trivial JPA projects now recognise Hibernate behind the abstraction
- ❖ Many non-trivial Hibernate projects now recognise SQL behind the abstraction
- ❖ And by the way: Most JSF projects now recognise JS and HTTP (REST)

Non-trivial ORM

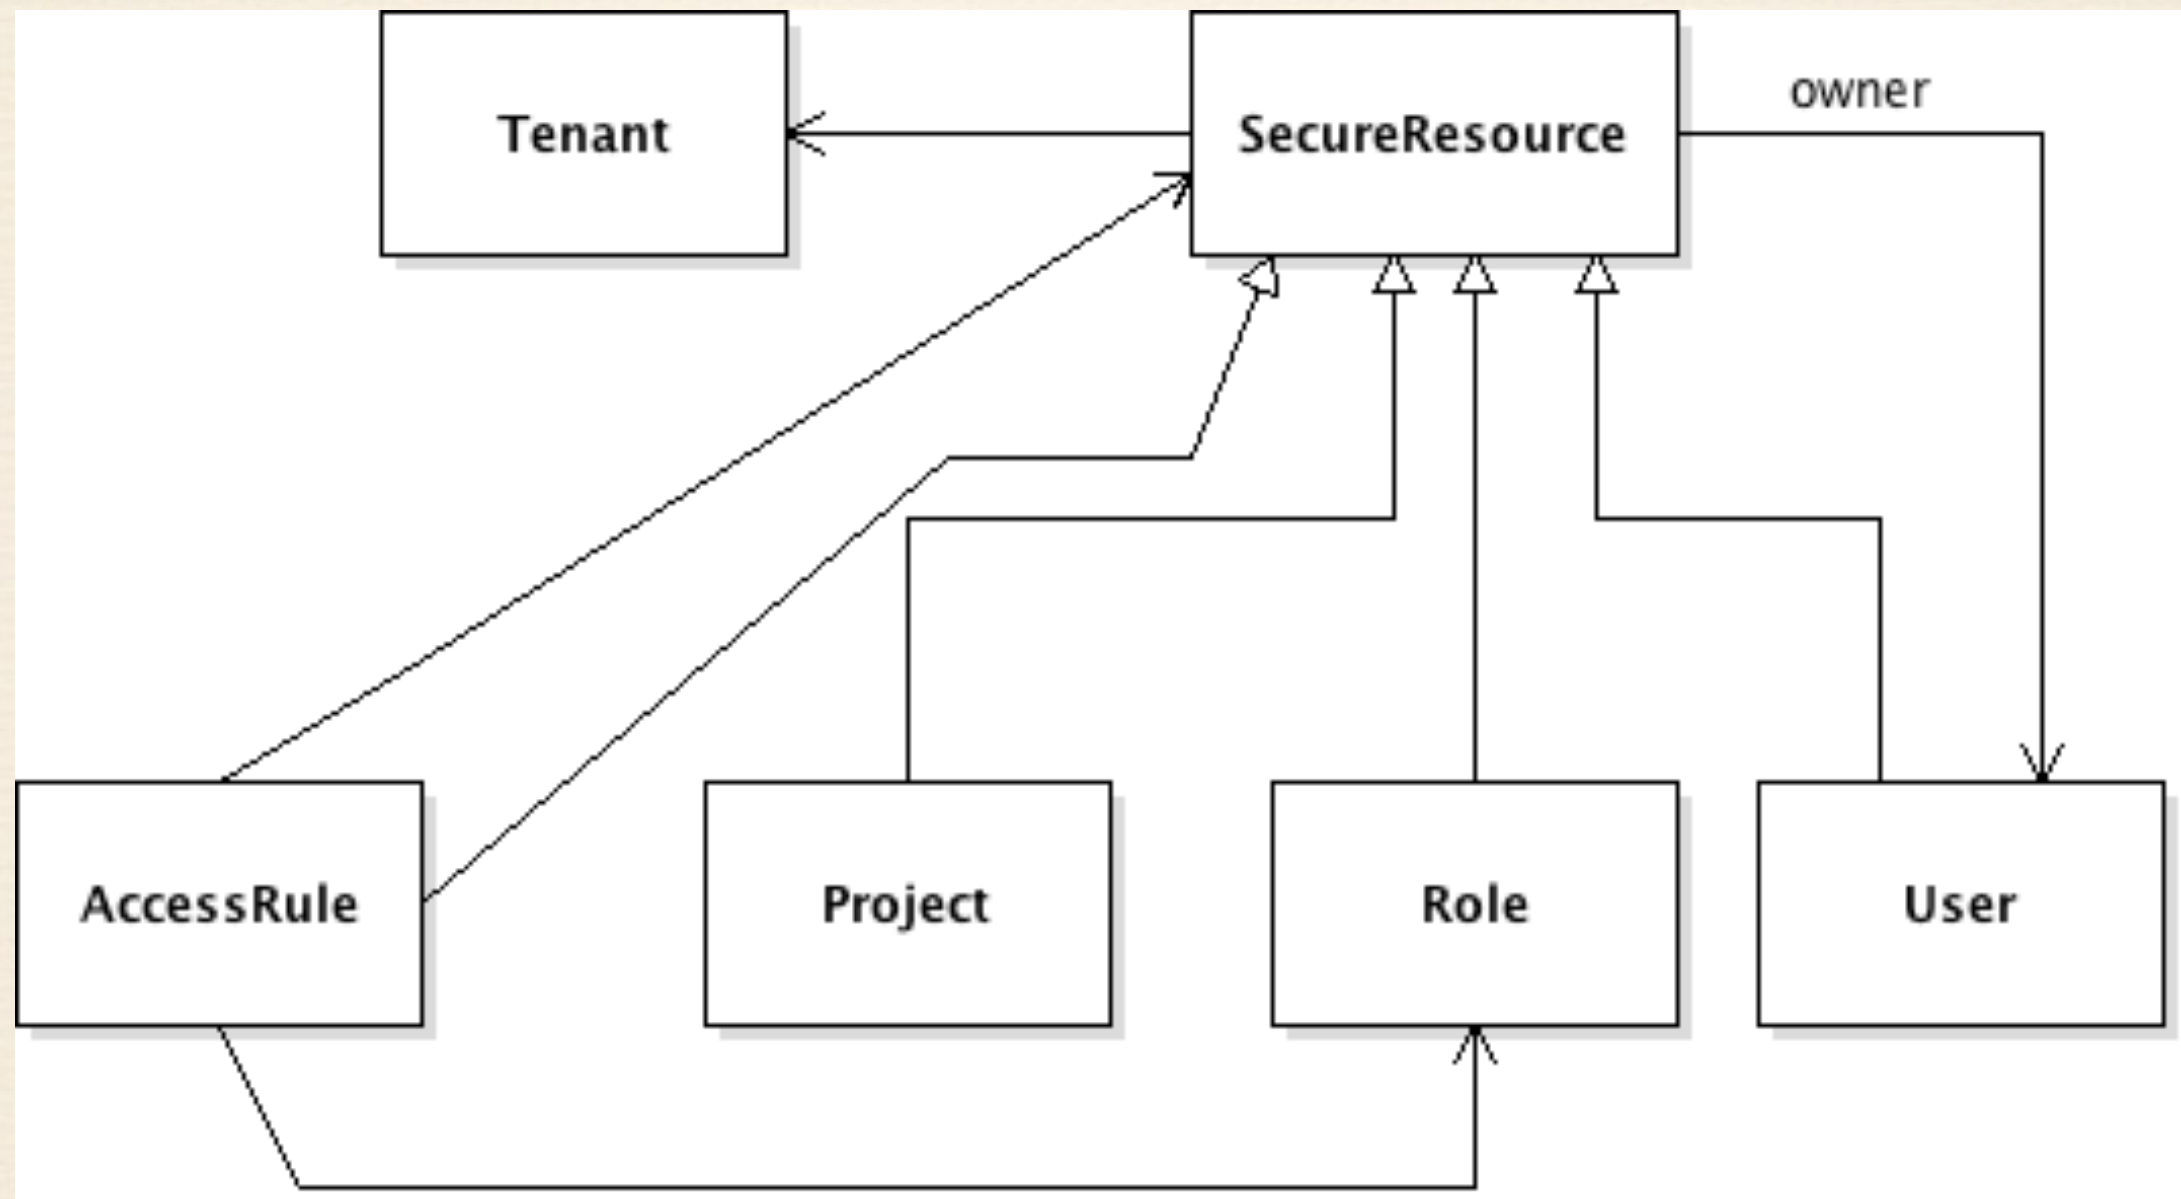
- ❖ Do I need ORM mapping, query abstraction?
- ❖ How smart do I need ORM to be?
- ❖ “There’s already one perfect DSL to access an RDBMS, it’s called SQL and it’s truly well-established - so why invent another?”
- ❖ Some current projects avoid ORM

ORM Gone Bad:

An example

- ❖ A SecureResource “knows” its owner and tenant
- ❖ A Project, a Role, a User and an AccessRule are all SecureResources
- ❖ An AccessRule associates Roles with SecureResources

Simple domain model



In the exercises

- ❖ Hibernate produces massive SQL Statements
- ❖ Not all Inheritance strategies work
- ❖

If the price is high,...

- ❖ Do I need referential Integrity via DB?
- ❖ Do I need to map inheritance?
- ❖ Is the domain model too tightly coupled?
- ❖ Do I need ACID Transactions across entities and multiple statements?
- ❖ Cut your domain model into reasonable parts

Careful Design your data model

- ❖ The domain model MUST consider the implementation details.
- ❖ Decouple aspects from the domain model
- ❖ Favour reference over inheritance
- ❖ consider “Embedded” mapping to avoid inheritance

More Data Model Design Aspects

- ❖ Do we really need normalisation? Disks are cheap.
- ❖ Identify clusters in your domain model, clusters may then become documents
- ❖ Loose coupling may improve performance
- ❖ Lazy-load hell can be avoided.

New kid on the block: CQRS

- ❖ Command-Query-Responsibility Segregation
- ❖ Separate reading to and writing from storage
- ❖ Optimise reading and writing independently
- ❖ Consider for very large and fast data requirements
- ❖ see: <http://martinfowler.com/bliki/CQRS.html>

Not only SQL

- ❖ What else?
- ❖ MongoDB, Redis, Cassandra, Hazelcast, Neo4J,...
- ❖ No more ACID, but CAP
- ❖ Remember: Architecture decisions are trade-off decisions - you may not get it all
- ❖ Cost of failure vs. cost of prevention

NoSQL? No, SQL!

- ❖ Consider pros and cons, features and acceptance
- ❖ What exactly is it that you need?
- ❖ Why not both?

Selection Criteria

- ❖ What part of ACID/CAP do I need? Is it provided?
- ❖ Which performance requirements do I expect mid-term? Size, latency in R/W?
- ❖ Is the product well-established and well-supported?
- ❖ Can the operations dept support the runtime characteristics?

Design to cost

- ❖ Why use transactions?
- ❖ What's the price of an "accident"?
- ❖ What's the likelihood of an "accident" to happen
- ❖ What's the price of preventing the accident?

One Alternative: MongoDB

- ❖ High performance db journaling
- ❖ Flexible access patterns
- ❖ Document-style database
- ❖ Secondary indexes, huge tables, simple queries, extremely fast and eventually persistent.

Approaching Swiss Trusted Poker

- ❖ No typical session state, thus RESTful.
- ❖ Multiple players share state: Tables and Games
 - ❖ Consider Actor-Based implementation or alike.
- ❖ Infrastructure like a poker table is held in MySQL
- ❖ Game state is held in Cache representing the game
- ❖ All messages are written to a persistent journal (MongoDB)
- ❖ The messages can recreate the state in case of failure.

In The Exercises:

Exploring Spring Data MongoDB

- ❖ Needs mongod installed on localhost
- ❖ Cool Abstraction
- ❖ Hiding the Persistence layer behind the repository interface.
- ❖ class: MongoPersonRepoTests
- ❖ Cool: PersonRepo is technology-agnostic

In the Exercises:

Exploring ORM Monster statements

- ❖ Some domain models are not easy to map
- ❖ Class InheritanceTests
- ❖ Creates massive statements
- ❖ Compare to class InheritanceTests2
- ❖ Domain model simplification solves the problem.

In the Exercises: Exploring Caching Aspects

- ❖ Caching is a typical cross-cutting concern
- ❖ Should be implemented as an aspect
- ❖ Class: CacheTests
- ❖ Pay attention to the eviction strategies
- ❖ Hint: search for and uncomment “show_sql”

Tonight's main course:

- ❖ Write a class that takes any
public Author findBy<Field> (String name)
interface and returns a JOOQ-based
implementation of it.